

## OS-5

1] Describe process creation with algorithm.

① Process creation :-

① The only way to create a new process in unix is to use the fork system calls. The process which calls fork is called the parent process & the newly created process is called child process.

② pid = fork();

③ pid for parent process is the process ID of child process. And pid for child process is 0. The process 0 is the only process which is not created via fork.

④ Steps followed by kernel for fork :-

i) It creates a new entry in process table  
ii) It assigns a unique ID to the newly created process.

iii) It makes logical copy of regions of parent process. If a regions can be shared, only its reference count is incremented instead of making a physical copy of region.

iv) The reference counts of file table entries & inodes of process are increased.

v) It turned the child process ID to parent & 0 to the child.

Algorithm :-

M	T	W	T	F	S	S
Page No.						
Date:					VOLUNTA	

Algorithm: fork

Input: none

Output: to parent process, child PID number  
to child process, 0

{

check for available kernel resources;  
get free proc table slot; unique PID number;  
check that user not running too many  
processes;

mark child state "being created";

copy data from parent proc table slot  
to new child slot;

increment counts on current directory  
inode and changed root (if applicable);

increment open file counts in file table;  
make copy of parent context

(u area, text, data, stack) in memory;  
push dummy system level context layer  
onto child system level context;

// dummy context contains data

// allowing child process to

// recognize itself & start running

// From here when scheduled

if (executing process is parent process)

{ if (executing process is parent process)

Change child state to "ready to run";  
return (child ID); // from system to user

? else // executing process is child process

Initialize u-area Timing fields;  
return (0);

Q.2] Explain algorithm for recognizing signals & handling signals.

- ① signals inform processes of the occurrence of asynchronous events.
- ② Handling signals:-
  - i) The kernel handles signals in context of process that receives them so a process must run to handle signals.
  - ii) old function = signal(signum, function);  
@ where signum is the signal number → the process is specifying the action for
  - b) function is the address of user function the process wants to invoke on receipt of signal and
  - c) return value oldfunction was the value of function in most recently specified call to signal for signum.
  - d) The process can pass the values 1 or 0 instead of function address:
    - The process will ignore future occurrence of signal if the parameter value is 1 &
    - exit in kernel on receipt of signal if its value is 0.
  - e) The u-area contains an array of signal-handlers fields, one for each signals defined in system. The kernel stores the address

Page 63		VGA
DATE		

of the user function in field that corresponds to signal number.

③ Alg :-

Algorithm: psig

Input : none

Output: none

{

    get signal number set in process table entry;  
 clear signal number in process table entry;  
 if (user had called signal call to ignore  
 this signal)

        return ;

//done

    if (user specified function to handle  
 the .signal)

{

        get user virtual address of signal catcher  
 stored in U area;

//the next statement has undesirable  
 side-effects

        clear U area entry that stored address  
 of signal catcher;

        modify user level context;

            artificially create user stack frame to  
 mimic call to signal catcher function;

        modify system level context;

            write address of signal catcher into  
 program counter field of user saved  
 register context;

        return ;

? {

    if (signal is type that system should dump  
 core image of process)

{

create file name "core" in current directory;  
write contents of user level context to  
file "core";

?

invoke exit algorithm immediately ;

?

④ IF signal handling function is set to  
default value, the kernel will dump a "core"  
image of process for certain types of signals  
before exiting.

⑤ If process decides to ignore, the signal  
field is not reset & the process will  
continue ignoring the signal.

⑥

Q. e] Explain exit system call with algorithm.

- ① Processes on UNIX system exit by executing the exit system call.
- ② When a process exits, it enters the zombie state, releases all of its resources and its context except for its process table entry.
- ③ exit (status);
  - where status is the exit code returned to parent.
  - The process may call exit explicitly, but the startup routine in C calls exit after the main function returns.
- ④ Algorithm: exit
  - Input: return code for parent process
  - Output: none.

{

  - ignore all signals;
  - if (process group leader with associated control terminal)

}

  - send hangup signal to all members of the process group;
  - reset process group for all members to 0;

{

  - close all open files (internal version of algorithm close);
  - release current directory (algorithm: input);
  - release current (changed) root, if exists (algorithm: input);

free regions, memory associated with process (algorithm : freeeg);  
write accounting record;  
make process state zombie;  
assign parent process ID for all child processes to be init process (1);  
if any children were zombie, send death of child signal to init;  
send death of child signal to parent process;  
context switch;

?

Q.4] describe awaiting process termination.

- ① A process can synchronize its execution with the termination of child process by executing the wait system call.
- ②  $\text{P pid} = \text{Wait(stat\_addr)}$ ;  
— where pid is process ID of zombie child, and stat\_addr is the address in user space of an integer that will contain the exit status code of child.
- ③ The kernel does not "long jump" from sleep, but returns wait.
- ④ The kernel will restart the wait loop, find zombi child at least one is guaranteed to exist, release the child's process table slot and return from wait system call.

⑤ If the process catches "death of child" signals, the kernel arranges to call the user signal-handling routine, as it does for other signals.

⑥ If the process ignores "death of child" signals, the kernel restarts the wait loop, frees the process table slot of zombie children and searches for more children.

Q. 5] Explain algorithm for booting the system.

→ Algorithm

\* Algorithm : start

\* Input : none

\* Output : none

\*/

```
{
    initialize all kernel data structures;
    pseudo-mount of root;
    hand-craft environment of process 0;
    fork process 1;
}

//process 1 here:
allocate region;
attach region to init address space;
grow region to accommodate code about
to copy in;
copy code from kernel space to init
use space to exec init;
change mode: return from kernel to
user mode;
/* init never gets here -- as result of
above change mode,
* init exec's /etc/init & becomes a
"normal" user process
* with respect to invocation of system
calls
```

\*/

```
? // proc 0 continues here
for kernel processes;
/* process 0 invokes the swapper to manage
the allocation of process address space
to main memory & swap devices.
```

M	T	W	T	F	S
Deja Vu					
0215					YOLVA

- \* This is an infinite loop; process 0 usually sleeps in the loop unless there is work for it to do.
- \* execute code for swapper algorithm;

p. 6 write short note

### a User id of process

- ① There are two user ID associated with a process, the real user ID & the effective user ID or setuid (set user ID)
- ② The real user ID identifies the user who is responsible for running process.
- ③ The effective User ID is used to assign ownership of newly created files to check file access permissions and to check permission to send signals to processes via kill system call.

### ④ setuid (uid);

— where uid is new user ID, and its result depends on current value of effective user ID.

— If the effective user ID of calling process is superuser, the kernel resets the real & effective user ID fields in process table & u-area to uid.

— If its not superuser, the kernel resets the effective user ID in the u-area

Project No.	Name	Date
10001	TOONI	

to uid, if uid has the value of real user ID or if it has value of the same user ID. otherwise, the system call return an error. Generally, a process inherits its real & effective user ID from its parent. during the fork system calls maintains their values across exec system calls.

### b) shell :

- ① A shell provides you with an interface to Unix system.
- ② It gathers input from you & executes programs based on that input.
- ③ When a program finishes executing, it displays that program's output.
- ④ Shell is an environment in which we can run our commands, programs & shell scripts
- ⑤ If the shell recognizes the input string as a built-in command, it executes the command internally without creating new process; otherwise, it assumes the command is the name of an executable file.

for ex. → ls | wc -l

M	T	W	T	F	S	S
Page No.						
Date						

## Tutorial No.7

Q. 1] Explain algorithm for process scheduling.

① Algorithm: schedule-process

    input : none

    output : none

    { while (no process picked to execute)

        for (every process on run queue)  
            pick highest priority that is loaded  
            in memory ;

        if (no process eligible to execute)  
            idle the machine ;

    /\* interrupt takes machine out of idle state \*/

    remove chosen process from run queue ;  
    switch context to that of chosen process,  
    resume its execution ;

    }

② ① The kernel executes the algorithm to  
    schedule a process.

② selects the highest priority process from  
    those in the states "ready to run",  
    loaded in memory and "preempted".

③ It does not select a process that is not  
    in main memory.

④ If several processes tie for highest  
    priority, the kernel picks the one that has  
    been "ready to run" for longest time.

Q. 2]

Page No.	Date	Page No.

⑤ If there are no processes eligible for execution, the processor idles until the next interrupt.

Q. 2 Describe scheduling parameters used in process scheduling.

→ ① Each process table entry contains a priority field for process scheduling.

② The priority of process in user mode is a function of its recent CPU usage.

lower priority if a process has recently used the CPU.

③ The range of process priorities are partitioned into two classes:

User priorities and kernel priorities.

④ each class contains several priority values. each priority has queue of processes logically associated with it.

⑤ User level priorities are below a threshold value and kernel-level priorities are above the threshold value.

⑥ Processes with user level priorities were preempted on their return from the kernel to user mode.

⑦ Processes with kernel-level priorities achieved them in the sleep algorithm.

⑧ Kernel-level priorities are further subdivided:

i) Processes with low kernel priority wake up on receipt of signal, but processes with

M	T	W	T	F	S	S
Page No.						
Date:	YOUVA					

high kernel priority continue to sleep.

i) The priorities called "swappet", "waiting for disk I/O", "Waiting for buffer", and "Waiting for inode" are high, noninterruptible system priorities.

Q. 3 Explain various system call for time.

M	T	W	T	F	S	S
Page No.:						
Date:					VOUVA	

4) Describe fun. of clock interrupt handler.

① The function of clock interrupt handler are to

- restart the clock
- schedule invocation of internal kernel functions based on internal timers.
- provide execution profiling capability for the kernel and for user processes
- gather system & process accounting statistics
- keep track of time
- send alarm signals to processes on request
- periodically wake up the swapper process
- control process scheduling

② Some operations are done every clock interrupt, whereas others are done after several clock ticks.

③ At every clock interrupt, the clock handler decrements the time field of first entry.

④ If the time field of first entry is less than or equal to 0 then, the specified function should be invoked.

- the clock handler does not invoke the fun directly.

⑤ The clock handler schedules the fun by causing a "software interrupt".

- because software interrupts are at lower priority level than other interrupts they are blocked until kernel finishes



handling all other interrupts.

— clock interrupts could occur until  
the software interrupt occurs and,  
therefore, the time field of first entry can  
have a negative value.